


# *BYOR: BRING YOUR OWN ROOTKIT*

Nischay Hegde  
Malware Detection Researcher  
SentinelOne





# Agenda

- Introduction
- Why?
- Goals of this talk
- How?
- Make it
- Detect it



# \$whoami

- Malware Detection Researcher @ SentinelOne
- Ex-Security Researcher @ Uptycs
- Linux Security nerd
- Contributes to random FOSS projects when they have issues that aggravate me
- Helping fix bugs and push for a FOSS messenger (Matrix).

# What is a Rootkit?

- A rootkit is a piece of software that hides its existence and allows other (unprivileged) programs to execute as root.
- Essentially, a very specific type of backdoor
- Usually installed by escalating privileges and using various vulnerabilities.



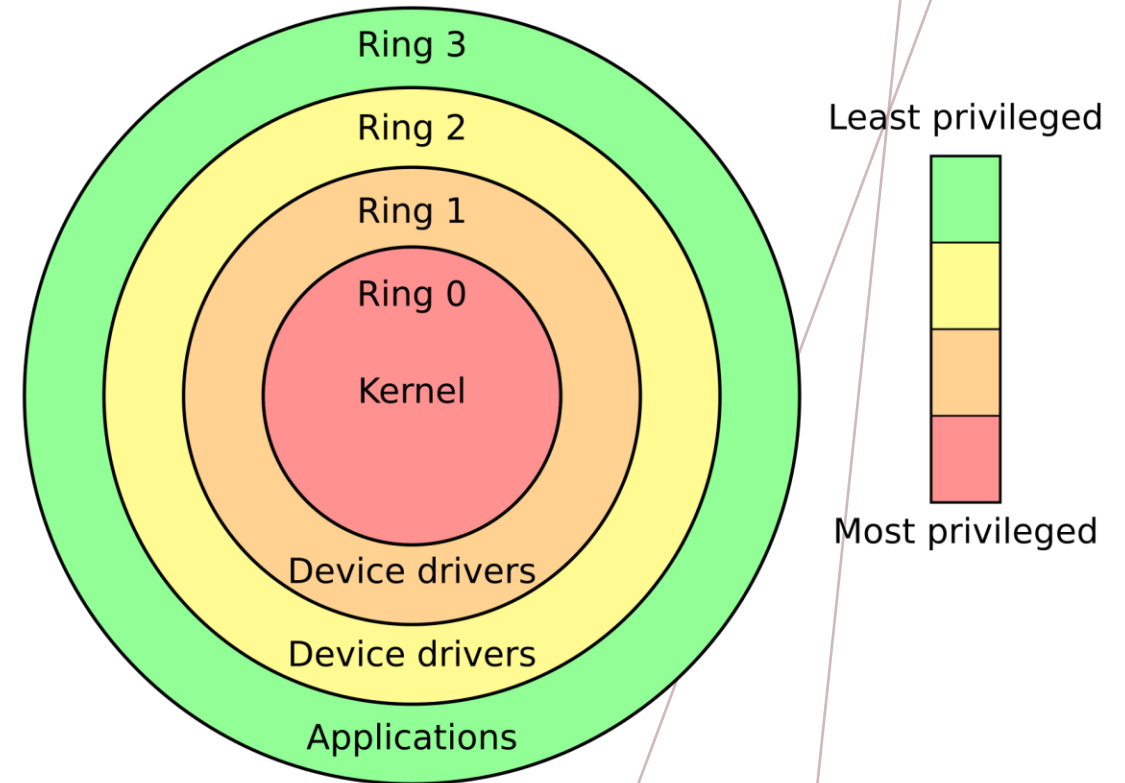
# Why would you make your own?

- Because it's interesting



# How?

- Rootkits are simply a subset of Linux Kernel Modules.
- LKMs are programs that get loaded as part of the kernel and run code in the kernel level.
- Can be listed with ``lsmod``, inserted with ``modprobe``, and removed by using ``rmmod``.
- Therefore, we need to at least defeat these three to make it “invisible”.



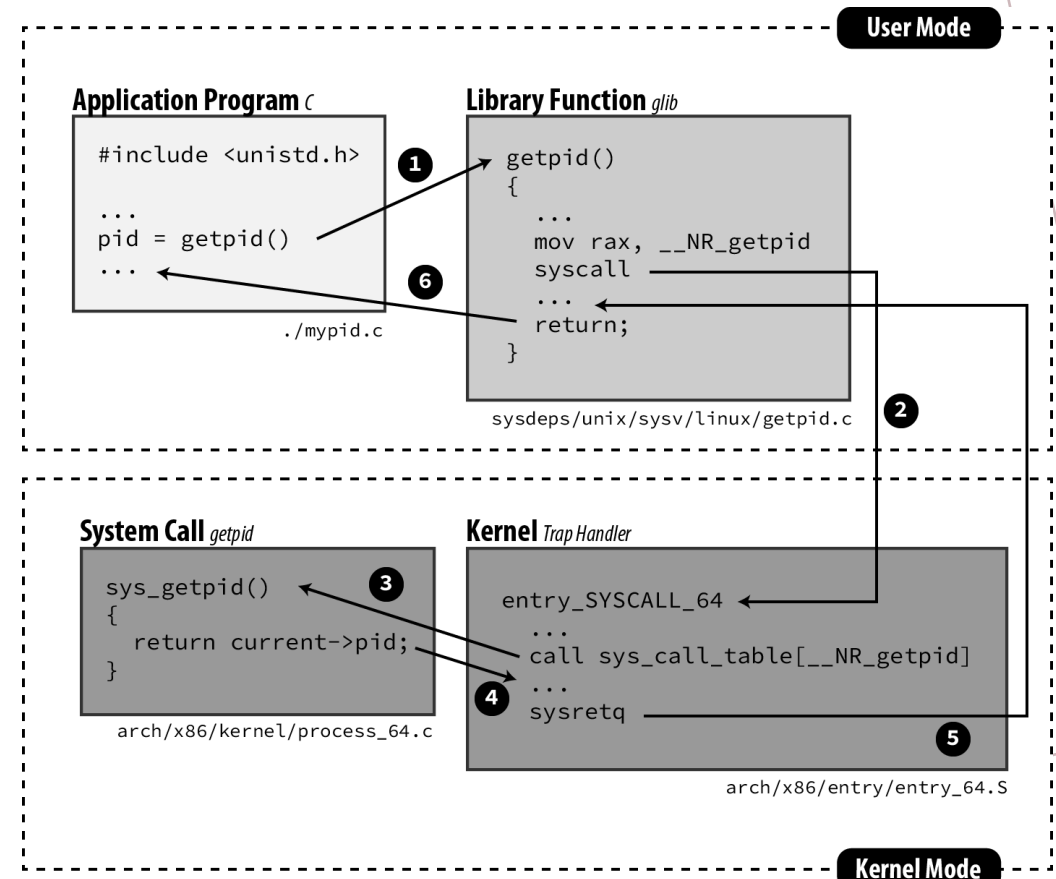
By Hertzprung at English Wikipedia, CC BY-SA 3.0,  
<https://commons.wikimedia.org/w/index.php?curid=8950144>

# Goals

- Write a kernel module that becomes invisible once you load it.
- Use it to do rootkit activities like giving root access to shells that do not have it.

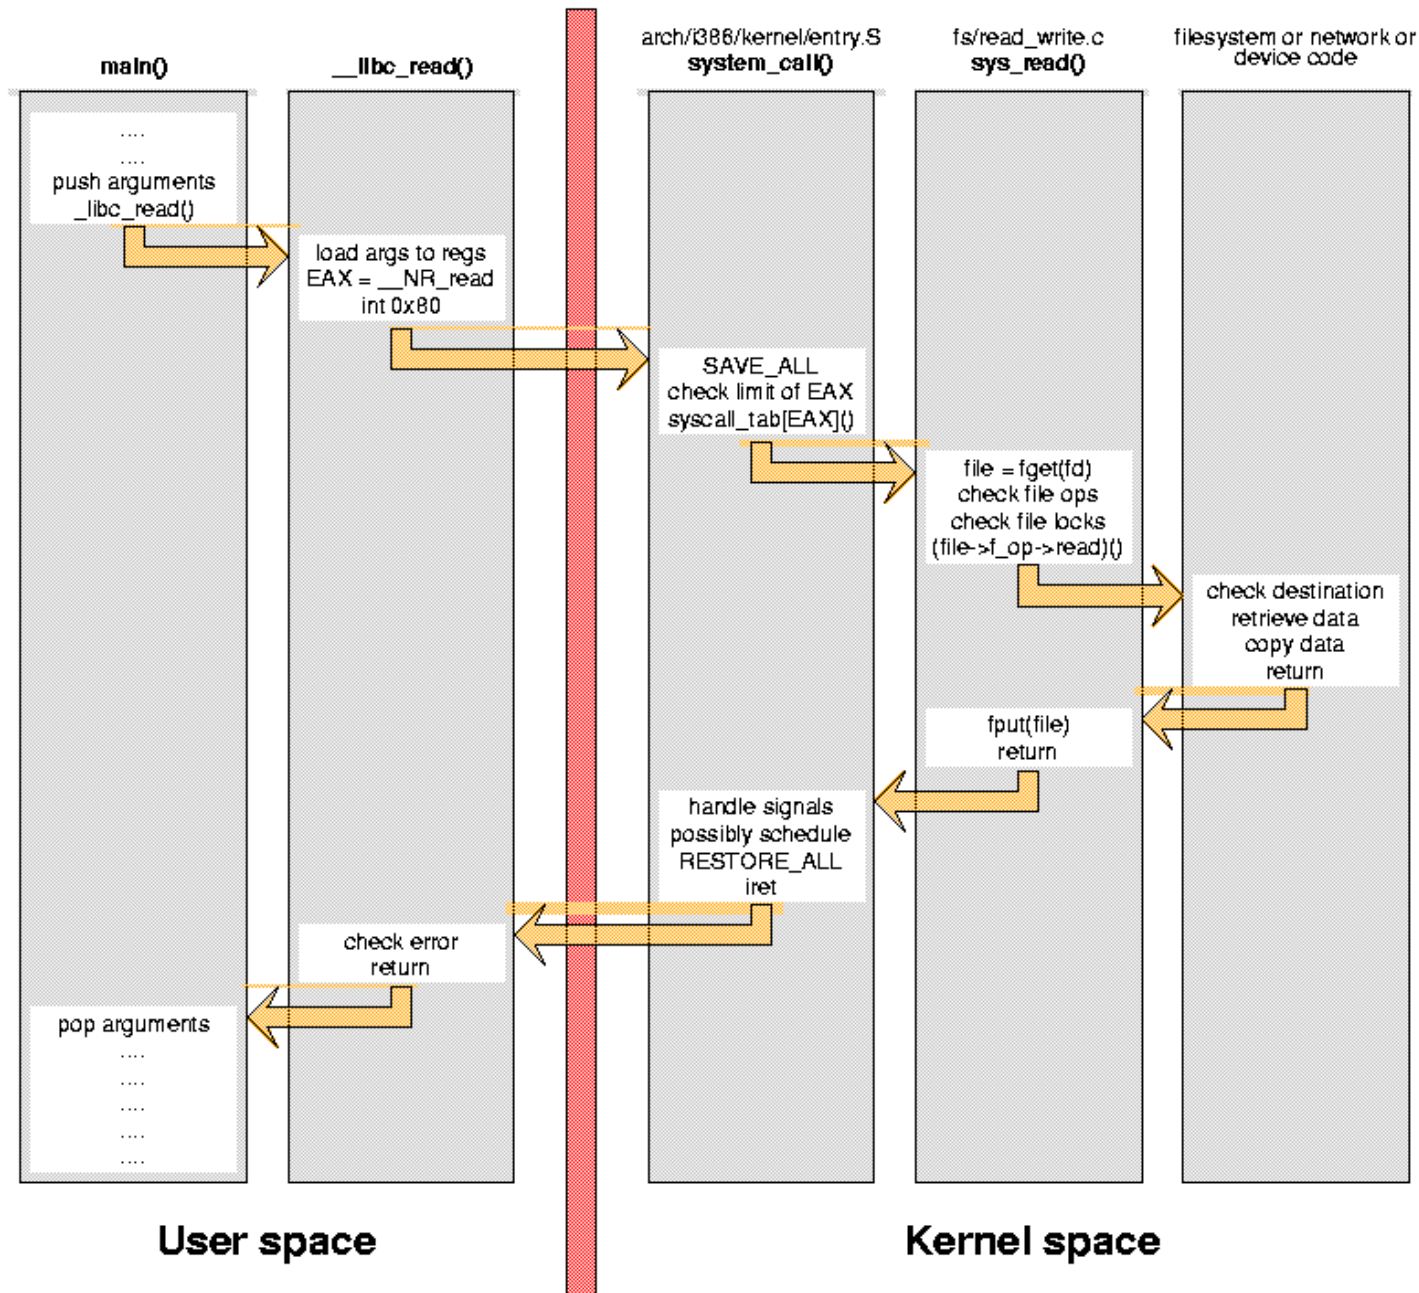
# Aside: How do syscalls work?

- Syscalls are an interface between the usermode and kernel-mode.
- When the syscall is called, the OS goes to the syscall table to find addresses associated with the syscall.
- In the syscall table, we redirect it to a function that does the required work.
- Therefore, our rootkit would have to hijack this process.





# Aside: How do syscalls work?

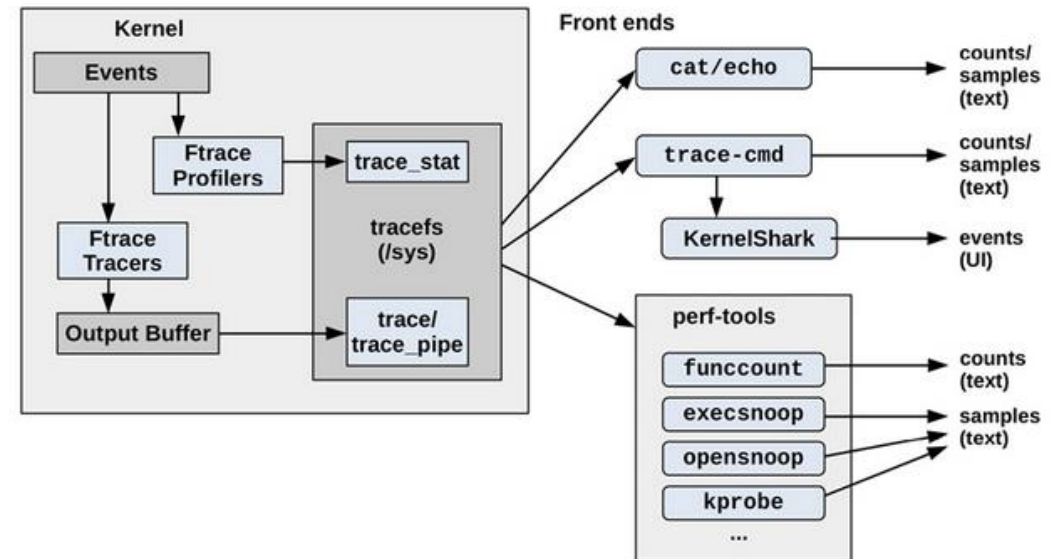


# How would the rootkit work?

- Our hypothetical rootkit would need to modify (or “hook”) into certain syscalls, change their behaviour and then make it do the work that we need it to do.
- Two popular types of hooking:
  - Hook the syscall table with either the new function or an inline assembly that calls the new function
  - ftrace
  - eBPF
- We will be using ftrace because hooking syscall table is mostly useless [now](#). (Linux 6.9+)

# What is Ftrace?

- Linux kernel tracing mechanism
- Built into kernel, no extra installation necessary
- Can hook into various probes (i.e. places) in the kernel and outside (they're called kprobes, uprobes, tracepoints etc)



# Aside: Linux kernel tracers

- Try to guess how many tracers exist.

It's 9!

- ftrace
- perf\_events
- eBPF
- SystemTap
- LTTng
- ktap
- dtrace4linux
- OL DTrace
- sysdig



**THIS IS FINE**



# Libraries that make it (slightly) easier

- <https://github.com/milabs/khook>
- <https://github.com/WeiJiLab/kernel-inline-hook-framework>
- Both are outdated (Linux 4.x mostly)





DEMO



# Detection

- We will be using eBPF to detect this attack.
- Another kernel observability tool.
- Increasingly being used by EDRs because it's less code on the kernel than kernel modules.
- Has a high level language to write code in (bpftrace) as well as a lower-level libraries for C and Python (BCC).



# References

- <https://www.brendangregg.com/blog/2015-07-08/choosing-a-linux-tracer.html>
- [https://github.com/xcellerator/linux\\_kernel\\_hacking](https://github.com/xcellerator/linux_kernel_hacking)
  - XCellerator's blog is XCellent as well
- <https://github.com/iovisor/bcc/blob/master/examples/tracing/stacksnoop.py> - stacksnoop



# *THANK YOU*

Nischay Hegde

mailto: [contact@nischay.me](mailto:contact@nischay.me)

Mastodon: [@thatloststudent@infosec.exchange](https://infosec.exchange/@thatloststudent)

Twitter: [@thatloststudent](https://twitter.com/thatloststudent)